

SKANDA VAIDYANATH

Founding AI Researcher, Yutori

@ s vaidyan@stanford.edu

[skandavaidyanath.github.io](https://github.com/skandavaidyanath)

in [skanda-vaidyanath](https://www.linkedin.com/in/skanda-vaidyanath)

[skandavaidyanath](https://www.kaggle.com/skandavaidyanath)

[Skanda Vaidyanath](https://www.instagram.com/skanda.vaidyanath)

EDUCATION

M.S. Computer Science A.I. Track

Stanford University

📅 Sep 2021 – Jun 2023

- CGPA: 4.06/4

B.E. (Hons.) Computer Science with a Minor in Data Science

BITS Pilani, Hyderabad Campus

📅 Jun 2016 – May 2020

- Class valedictorian.
- CGPA: 9.93/10, Major CGPA (only CS courses): 10.00/10

RECENT EXPERIENCE

Founding AI Researcher

Yutori

📅 Feb 2025 – Present

📍 San Francisco, California

- RL post-training for the [Navigator n1](#), [n1.5](#), and [n2](#) model series (GRPO variants, value function methods, and [SDPO](#)).
- [Navigator n2](#) is a 27B computer-use model achieving **SOTA** results on four of five evaluated benchmarks.
- Built the agentic harness and backend for [Scouts](#) and [Delegate](#).

Staff Research Scientist

Riot Games AI Accelerator

📅 Jul 2023 – Feb 2025

📍 Redwood City, California

- Trained self-play policies for two-player zero-sum partially observable stochastic games, achieving professional-level performance.
- Developed game-theoretic extensions to [asynchronous PPO](#) based on [R-NaD](#) and [magnetic mirror descent \(MMD\)](#).
- Scaled self-play using large-scale distributed reinforcement learning.

Research Engineer Intern

Google DeepMind

📅 Jun 2022 – Sep 2022

📍 Mountain View, California

- Co-developed [PushWorld](#), a benchmark of 200+ puzzles for manipulation planning with movable obstacles and tools [\[paper\]](#).
- Evaluated classical planners and model-free RL agents; introduced a planning heuristic 40× faster than the strongest baseline.

Research Intern

Microsoft Research

📅 Dec 2020 – Jul 2021

📍 Bangalore, India

- Built [Jigsaw](#), a multi-modal system that synthesizes Pandas code from natural language and I/O examples using GPT-3 and Codex [\[paper\]](#).
- Developed program-analysis and synthesis post-processing to repair reference, argument, and semantic errors and learn from user feedback.

AWARDS, HONORS



BITS Hyderabad Merit Scholarship for finishing in the top 1% of my graduating class every semester



IUSSTF-Viterbi Scholar 2019. I was **one of fifteen students** chosen from India for the programme.



Max Planck Institute for Informatics 2019 Research Scholar Fellowship

SELECT PUBLICATIONS

- Akash Velu^{*}, **Skanda Vaidyanath^{*}**, Dilip Arumugam. Hindsight-DICE: Stable Credit Assignment for Deep Reinforcement Learning. *arXiv preprint*, 2023. [\[pdf\]](#)
- Divyansh Garg^{*}, **Skanda Vaidyanath^{*}**, Kuno Kim, Jiaming Song, Stefano Ermon. LISA: Learning Interpretable Skill Abstractions from Language. *NeurIPS 2022* [\[pdf\]](#) [\[website\]](#)
- Naman Jain, **Skanda Vaidyanath**, Arun Iyer, Nagarajan Natarajan, Suresh Parthasarathy, Sriram Rajamani, Rahul Sharma. Jigsaw: Large Language Models meet Program Synthesis. *ICSE 2022* [\[pdf\]](#)[\[blog\]](#)

ACTIVITIES

- Course Assistant, Deep Learning: CS230 Autumn 2022, Spring 2023
- Course Assistant, Artificial Intelligence: CS221 Spring 2022
- Course Assistant, Reinforcement Learning: CS234 Winter 2022, Winter 2023
- Research With Impact, SGSI 2021, Stanford University [\[link\]](#)
- Leadership Labs, SGSI 2020, Stanford University [\[link\]](#)

ML FRAMEWORKS

[Miles](#)

[verl](#)

[LLaMA Factory](#)

[Transformers](#)

[vLLM](#)

[SGLang](#)

SELECT COURSEWORK

[Probabilistic Graphical Models](#)

[Convex Optimization](#)

[Reinforcement Learning](#)

[Meta-Learning](#)

[Graph Neural Networks](#)

[Randomized Algorithms](#)